

**a5/cuda\_kmeans\_naive.cu**

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 #include "kmeans.h"
5 #include "alloc.h"
6 #include "error.h"
7
8 #ifdef __CUDACC__
9 inline void checkCuda(cudaError_t e) {
10     if (e != cudaSuccess) {
11         // cudaGetErrorString() isn't always very helpful. Look up the error
12         // number in the cudaError enum in driver_types.h in the CUDA includes
13         // directory for a better explanation.
14         error("CUDA Error %d: %s\n", e, cudaGetErrorString(e));
15     }
16 }
17
18 inline void checkLastCudaError() {
19     checkCuda(cudaGetLastError());
20 }
21 #endif
22
23 __device__ int get_tid() {
24     return blockIdx.x * blockDim.x + threadIdx.x;
25 }
26
27 /* square of Euclid distance between two multi-dimensional points */
28 __host__ __device__ inline static
29 double euclid_dist_2(int numCoords,
30                      int numObjs,
31                      int numClusters,
32                      double *objects,      // [numObjs][numCoords]
33                      double *clusters,     // [numClusters][numCoords]
34                      int objectId,
35                      int clusterId) {
36     int i;
37     double ans = 0.0;
38
39     /* TODO: Calculate the euclid_dist of elem=objectId of objects from elem=clusterId from
clusters*/
40     for (i = 0; i < numCoords; i++) {
41         double objectVal = objects[objectId * numCoords + i];
42         double clusterVal = clusters[clusterId * numCoords + i];
43
44         double diff = objectVal - clusterVal;
45         ans += diff * diff;
46     }
47
48     return (ans);
49 }
50
51 __global__ static
```

```
52 void find_nearest_cluster(int numCoords,
53                           int numObjs,
54                           int numClusters,
55                           double *objects,           // [numObjs][numCoords]
56                           double *deviceClusters,    // [numClusters][numCoords]
57                           int *deviceMembership,     // [numObjs]
58                           double *devdelta) {
59
60     /* Get the global ID of the thread. */
61     int tid = get_tid();
62
63     if (tid < numObjs) {
64         int index, i;
65         double dist, min_dist;
66
67         /* find the cluster id that has min distance to object */
68         index = 0;
69
70         min_dist = euclid_dist_2(numCoords, numObjs, numClusters,
71                               objects, deviceClusters,
72                               tid, index);
73
74         for (i = 1; i < numClusters; i++) {
75
76             dist = euclid_dist_2(numCoords, numObjs, numClusters,
77                               objects, deviceClusters,
78                               tid, i);
79             /* no need square root */
80             if (dist < min_dist) { /* find the min and its array index */
81                 min_dist = dist;
82                 index = i;
83             }
84         }
85
86         if (deviceMembership[tid] != index) {
87
88             atomicAdd(devdelta, 1.0);
89         }
90
91         /* assign the deviceMembership to object objectId */
92         deviceMembership[tid] = index;
93     }
94 }
95
96 //
97 // -----
98 // DATA LAYOUT
99 //
100 // objects      [numObjs][numCoords]
101 // clusters     [numClusters][numCoords]
102 // newClusters   [numClusters][numCoords]
103 // deviceObjects  [numObjs][numCoords]
104 // deviceClusters [numClusters][numCoords]
105 // -----
```

```

106 //                                                 */
107 /* return an array of cluster centers of size [numClusters][numCoords]      */
108 void kmeans_gpu(double *objects,          /* in: [numObjs][numCoords] */
109                  int numCoords,        /* no. features */
110                  int numObjs,         /* no. objects */
111                  int numClusters,      /* no. clusters */
112                  double threshold,     /* % objects change membership */
113                  long loop_threshold,   /* maximum number of iterations */
114                  int *membership,      /* out: [numObjs] */
115                  double *clusters,       /* out: [numClusters][numCoords] */
116                  int blockSize) {
117
118     double timing = wtime(), timing_internal, timer_min = 1e42, timer_max = 0;
119     double timing_gpu, timing_cpu, timing_transfers, transfers_time = 0.0, cpu_time = 0.0,
120     gpu_time = 0.0;
121
122     int loop_iterations = 0;
123     int i, j, index, loop = 0;
124     int *newClusterSize; /* [numClusters]: no. objects assigned in each
125                           new cluster */
126
127     double delta = 0, *dev_delta_ptr;           /* % of objects change their clusters */
128     double **newClusters = (double **) calloc_2d(numClusters, numCoords, sizeof(double));
129
130     double *deviceObjects;
131     double *deviceClusters;
132     int *deviceMembership;
133
134     printf("\n|-----Naive GPU Kmeans-----|\n\n");
135
136
137     /* initialize membership[] */
138     for (i = 0; i < numObjs; i++) membership[i] = -1;
139
140     /* need to initialize newClusterSize and newClusters[0] to all 0 */
141     newClusterSize = (int *) calloc(numClusters, sizeof(int));
142     assert(newClusterSize != NULL);
143
144     timing = wtime() - timing;
145     printf("t_alloc: %lf ms\n\n", 1000 * timing);
146     timing = wtime();
147
148     const unsigned int numThreadsPerClusterBlock = (numObjs > blockSize) ? blockSize :
149     numObjs;
150
151     const unsigned int numClusterBlocks = (numObjs + numThreadsPerClusterBlock - 1) /
152     numThreadsPerClusterBlock; /* TODO: Calculate Grid size, e.g. number of blocks. */
153
154     const unsigned int clusterBlockSharedDataSize = 0;
155
156     checkCuda(cudaMalloc(&deviceObjects, numObjs * numCoords * sizeof(double)));
157     checkCuda(cudaMalloc(&deviceClusters, numClusters * numCoords * sizeof(double)));
158     checkCuda(cudaMalloc(&deviceMembership, numObjs * sizeof(int)));
159     checkCuda(cudaMalloc(&dev_delta_ptr, sizeof(double)));
160
161     timing = wtime() - timing;
162     printf("t_alloc_gpu: %lf ms\n\n", 1000 * timing);
163     timing = wtime();
164
165

```

```
157     checkCuda(cudaMemcpy(deviceObjects, objects,
158                         numObjs * numCoords * sizeof(double), cudaMemcpyHostToDevice));
159     checkCuda(cudaMemcpy(deviceMembership, membership,
160                         numObjs * sizeof(int), cudaMemcpyHostToDevice));
161     timing = wtime() - timing;
162     printf("t_get_gpu: %lf ms\n\n", 1000 * timing);
163     timing = wtime();
164
165 do {
166     timing_internal = wtime();
167
168     /* GPU part: calculate new memberships */
169
170     timing_transfers = wtime();
171
172     checkCuda(cudaMemcpy(deviceClusters, clusters,
173                         numClusters * numCoords * sizeof(double),
174                         cudaMemcpyHostToDevice));
175
176     transfers_time += wtime() - timing_transfers;
177
178     checkCuda(cudaMemset(dev_delta_ptr, 0, sizeof(double)));
179
180     //printf("Launching find_nearest_cluster Kernel with grid_size = %d, block_size = %d,
181     //shared_mem = %d KB\n", numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDa-
182     taSize/1000);
183     timing_gpu = wtime();
184     find_nearest_cluster
185     <<< numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDataSize >>>
186         (numCoords, numObjs, numClusters,
187          deviceObjects, deviceClusters, deviceMembership, dev_delta_ptr);
188
189     cudaDeviceSynchronize();
190     checkLastCudaError();
191     gpu_time += wtime() - timing_gpu;
192     //printf("Kernels complete for itter %d, updating data in CPU\n", loop);
193
194     timing_transfers = wtime();
195
196     checkCuda(cudaMemcpy(membership, deviceMembership,
197                         numObjs * sizeof(int),
198                         cudaMemcpyDeviceToHost));
199
200     checkCuda(cudaMemcpy(&delta, dev_delta_ptr,
201                         sizeof(double),
202                         cudaMemcpyDeviceToHost));
203
203     transfers_time += wtime() - timing_transfers;
204
205     /* CPU part: Update cluster centers*/
206     timing_cpu = wtime();
207     for (i = 0; i < numObjs; i++) {
208         /* find the array index of nestest cluster center */
209         index = membership[i];
```

```

209
210     /* update new cluster centers : sum of objects located within */
211     newClusterSize[index]++;
212     for (j = 0; j < numCoords; j++)
213         newClusters[index][j] += objects[i * numCoords + j];
214     }
215
216     /* average the sum and replace old cluster centers with newClusters */
217     for (i = 0; i < numClusters; i++) {
218         for (j = 0; j < numCoords; j++) {
219             if (newClusterSize[i] > 0)
220                 clusters[i * numCoords + j] = newClusters[i][j] / newClusterSize[i];
221             newClusters[i][j] = 0.0; /* set back to 0 */
222         }
223         newClusterSize[i] = 0; /* set back to 0 */
224     }
225
226     delta /= numObjs;
227     //printf("delta is %f - ", delta);
228     loop++;
229     //printf("completed loop %d\n", loop);
230     cpu_time += wtime() - timing_cpu;
231
232     timing_internal = wtime() - timing_internal;
233     if (timing_internal < timer_min) timer_min = timing_internal;
234     if (timing_internal > timer_max) timer_max = timing_internal;
235 } while (delta > threshold && loop < loop_threshold);
236
237 timing = wtime() - timing;
238 printf("\nloops = %d : total = %lf ms\n\t-> t_loop_avg = %lf ms\n\t-> t_loop_min = %lf
ms\n\t-> t_loop_max = %lf ms\n\t"
239         "-> t_cpu_avg = %lf ms\n\t-> t_gpu_avg = %lf ms\n\t-> t_transfers_avg = %lf
ms\n-----|\n",
240         loop, 1000 * timing, 1000 * timing / loop, 1000 * timer_min, 1000 * timer_max,
241         1000 * cpu_time / loop, 1000 * gpu_time / loop, 1000 * transfers_time / loop);
242
243 char outfile_name[1024] = {0};
244 sprintf(outfile_name, "Execution_logs/silver1-V100_Sz-%lu_Coo-%d_C1-%d.csv",
245         numObjs * numCoords * sizeof(double) / (1024 * 1024), numCoords, numClusters);
246 FILE *fp = fopen(outfile_name, "a+");
247 if (!fp) error("Filename %s did not open successfully, no logging performed\n",
outfile_name);
248 fprintf(fp, "%s,%d,%lf,%lf,%lf\n", "Naive", blockSize, timing / loop, timer_min,
timer_max);
249 fclose(fp);
250 checkCuda(cudaFree(deviceObjects));
251 checkCuda(cudaFree(deviceClusters));
252 checkCuda(cudaFree(deviceMembership));
253
254 free(newClusters[0]);
255 free(newClusters);
256 free(newClusterSize);
257
258 return;

```

259 | }

260 |

261 |