

a5/cuda_kmeans_shared.cu

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 #include "kmeans.h"
5 #include "alloc.h"
6 #include "error.h"
7
8 #ifdef __CUDACC__
9 inline void checkCuda(cudaError_t e) {
10     if (e != cudaSuccess) {
11         // cudaGetErrorString() isn't always very helpful. Look up the error
12         // number in the cudaError enum in driver_types.h in the CUDA includes
13         // directory for a better explanation.
14         error("CUDA Error %d: %s\n", e, cudaGetErrorString(e));
15     }
16 }
17
18 inline void checkLastCudaError() {
19     checkCuda(cudaGetLastError());
20 }
21 #endif
22
23 __device__ int get_tid() {
24     return blockIdx.x * blockDim.x + threadIdx.x;
25 }
26
27 /* square of Euclid distance between two multi-dimensional points using column-base format */
28 __host__ __device__ inline static
29 double euclid_dist_2_transpose(int numCoords,
30                                int numObjs,
31                                int numClusters,
32                                double *objects,      // [numCoords][numObjs]
33                                double *clusters,     // [numCoords][numClusters]
34                                int objectId,
35                                int clusterId) {
36     int i;
37     double ans = 0.0;
38
39     /* TODO: Calculate the euclid_dist of elem=objectId of objects from elem=clusterId from
clusters, but for column-base format!!! */
40     for (i = 0; i < numCoords; i++) {
41         double objectVal = objects[i * numObjs + objectId];
42         double clusterVal = clusters[i * numClusters + clusterId];
43
44         double diff = objectVal - clusterVal;
45         ans += diff * diff;
46     }
47
48     return (ans);
49 }
50
```

```

51 __global__ static
52 void find_nearest_cluster(int numCoords,
53                           int numObjs,
54                           int numClusters,
55                           double *objects,           // [numCoords][numObjs]
56                           double *deviceClusters,    // [numCoords][numClusters]
57                           int *deviceMembership,     // [numObjs]
58                           double *devdelta) {
59     extern __shared__ double shmemClusters[];
60
61     // TODO: Copy deviceClusters to shmemClusters so they can be accessed faster.
62     int tid_in_block = threadIdx.x;      // Το ID του νήματος μέσα στο Block
63     int block_size = blockDim.x;         // Πόσα νήματα έχει το Block
64     int total_cluster_doubles = numClusters * numCoords; // Συνολικά νούμερα προς αντιγραφή
65
66     // Κάθε νήμα αντιγράφει όσα στοιχεία του αναλογούν (με βήμα block_size)
67     for (int k = tid_in_block; k < total_cluster_doubles; k += block_size) {
68         shmemClusters[k] = deviceClusters[k];
69     }
70
71     /* Συγχρονισμός (BARRIER) */
72
73     __syncthreads();
74
75     /* Get the global ID of the thread. */
76     int tid = get_tid();
77
78     /* TODO: Maybe something is missing here... should all threads run this? */
79     if (tid < numObjs) {
80         int index, i;
81         double dist, min_dist;
82
83         /* find the cluster id that has min distance to object */
84         index = 0;
85         /* TODO: call min_dist = euclid_dist_2(...) with correct objectId/clusterId using
clusters in shmem*/
86
87
88         min_dist = euclid_dist_2_transpose(numCoords, numObjs, numClusters,
89                                           objects, shmemClusters,
90                                           tid, index);
91
92         for (i = 1; i < numClusters; i++) {
93             dist = euclid_dist_2_transpose(numCoords, numObjs, numClusters,
94                                           objects, shmemClusters,
95                                           tid, i);
96
97             /* no need square root */
98             if (dist < min_dist) { /* find the min and its array index */
99                 min_dist = dist;
100                index = i;
101            }
102        }
103    }

```

```

104     if (deviceMembership[tid] != index) {
105         /* TODO: Maybe something is missing here... is this write safe? */
106         atomicAdd(devdelta, 1.0);
107     }
108
109     /* assign the deviceMembership to object objectId */
110     deviceMembership[tid] = index;
111 }
112 }
113
114 //
115 // -----
116 // DATA LAYOUT
117 //
118 // objects      [numObjs][numCoords]
119 // clusters     [numClusters][numCoords]
120 // dimObjects   [numCoords][numObjs]
121 // dimClusters  [numCoords][numClusters]
122 // newClusters  [numCoords][numClusters]
123 // deviceObjects [numCoords][numObjs]
124 // deviceClusters [numCoords][numClusters]
125 //
126 //
127 /* return an array of cluster centers of size [numClusters][numCoords]      */
128 void kmeans_gpu(double *objects,          /* in: [numObjs][numCoords] */
129                  int numCoords,    /* no. features */
130                  int numObjs,     /* no. objects */
131                  int numClusters, /* no. clusters */
132                  double threshold, /* % objects change membership */
133                  long loop_threshold, /* maximum number of iterations */
134                  int *membership, /* out: [numObjs] */
135                  double *clusters, /* out: [numClusters][numCoords] */
136                  int blockSize) {
137     double timing = wtime(), timing_internal, timer_min = 1e42, timer_max = 0;
138     double timing_gpu, timing_cpu, timing_transfers, transfers_time = 0.0, cpu_time = 0.0,
gpu_time = 0.0;
139     int loop_iterations = 0;
140     int i, j, index, loop = 0;
141     int *newClusterSize; /* [numClusters]: no. objects assigned in each
new cluster */
142     double delta = 0, *dev_delta_ptr; /* % of objects change their clusters */
143     /* TODO: Copy me from transpose version*/
144     double **dimObjects = (double **) calloc_2d(numCoords, numObjs, sizeof(double));
//calloc_2d(...)->[numCoords][numObjs]
145     double **dimClusters = (double **) calloc_2d(numCoords, numClusters, sizeof(double));
//calloc_2d(...)->[numCoords][numClusters]
146     double **newClusters = (double **) calloc_2d(numCoords, numClusters, sizeof(double));
//calloc_2d(...)->[numCoords][numClusters]
147
148     double *deviceObjects;
149     double *deviceClusters;
150     int *deviceMembership;
151
152     printf("\n|-----Shared GPU Kmeans-----|\n\n");

```

```

154
155     /* TODO: Copy me from transpose version*/
156     for (i=0 ; i < numObjs; i++){
157         for (j=0; j<numCoords; j++){
158             dimObjects[j][i]=objects[i*numCoords + j];
159         }
160     }
161
162     /* pick first numClusters elements of objects[] as initial cluster centers*/
163     for (i = 0; i < numCoords; i++) {
164         for (j = 0; j < numClusters; j++) {
165             dimClusters[i][j] = dimObjects[i][j];
166         }
167     }
168
169     /* initialize membership[] */
170     for (i = 0; i < numObjs; i++) membership[i] = -1;
171
172     /* need to initialize newClusterSize and newClusters[0] to all 0 */
173     newClusterSize = (int *) calloc(numClusters, sizeof(int));
174     assert(newClusterSize != NULL);
175
176     timing = wtime() - timing;
177     printf("t_alloc: %lf ms\n\n", 1000 * timing);
178     timing = wtime();
179     const unsigned int numThreadsPerClusterBlock = (numObjs > blockSize) ? blockSize :
180     numObjs;
181     const unsigned int numClusterBlocks = (numObjs + numThreadsPerClusterBlock - 1) /
182     numThreadsPerClusterBlock; /* TODO: Calculate Grid size, e.g. number of blocks. */
183
184     /* Define the shared memory needed per block.
185      - BEWARE: We can overrun our shared memory here if there are too many
186      clusters or too many coordinates!
187      - This can lead to occupancy problems or even inability to run.
188      - Your exercise implementation is not requested to account for that (e.g. always
189      assume deviceClusters fit in shmemClusters */
190     const unsigned int clusterBlockSharedDataSize = numClusters*numCoords*sizeof(double);
191
192     cudaDeviceProp deviceProp;
193     int deviceNum;
194     cudaGetDevice(&deviceNum);
195     cudaGetDeviceProperties(&deviceProp, deviceNum);
196
197     if (clusterBlockSharedDataSize > deviceProp.sharedMemPerBlock) {
198         error("Your CUDA hardware has insufficient block shared memory to hold all cluster
199         centroids\n");
200     }
201
202     checkCuda(cudaMalloc(&deviceObjects, numObjs * numCoords * sizeof(double)));
203     checkCuda(cudaMalloc(&deviceClusters, numClusters * numCoords * sizeof(double)));
204     checkCuda(cudaMalloc(&deviceMembership, numObjs * sizeof(int)));
205     checkCuda(cudaMalloc(&dev_delta_ptr, sizeof(double)));
206
207     timing = wtime() - timing;

```

```
204     printf("t_alloc_gpu: %lf ms\n\n", 1000 * timing);
205     timing = wtime();
206
207     checkCuda(cudaMemcpy(deviceObjects, dimObjects[0],
208                          numObjs * numCoords * sizeof(double), cudaMemcpyHostToDevice));
209     checkCuda(cudaMemcpy(deviceMembership, membership,
210                         numObjs * sizeof(int), cudaMemcpyHostToDevice));
211     timing = wtime() - timing;
212     printf("t_get_gpu: %lf ms\n\n", 1000 * timing);
213     timing = wtime();
214
215 do {
216     timing_internal = wtime();
217
218     /* GPU part: calculate new memberships */
219
220     timing_transfers = wtime();
221     // TODO: Copy clusters to deviceClusters
222     checkCuda(cudaMemcpy(deviceClusters, dimClusters[0],
223                          numClusters * numCoords * sizeof(double),
224                          cudaMemcpyHostToDevice));
225
226     transfers_time += wtime() - timing_transfers;
227
228     checkCuda(cudaMemset(dev_delta_ptr, 0, sizeof(double)));
229
230     timing_gpu = wtime();
231     //printf("Launching find_nearest_cluster Kernel with grid_size = %d, block_size = %d,
232     shared_mem = %d KB\n", numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDataSize/1000);
233     find_nearest_cluster
234     <<< numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDataSize >>>
235         (numCoords, numObjs, numClusters,
236          deviceObjects, deviceClusters, deviceMembership, dev_delta_ptr);
237
238     cudaDeviceSynchronize();
239     checkLastCudaError();
240     gpu_time += wtime() - timing_gpu;
241     //printf("Kernels complete for itter %d, updating data in CPU\n", loop);
242
243     timing_transfers = wtime();
244
245     checkCuda(cudaMemcpy(membership, deviceMembership,
246                          numObjs * sizeof(int),
247                          cudaMemcpyDeviceToHost));
248
249     checkCuda(cudaMemcpy(&delta, dev_delta_ptr,
250                         sizeof(double),
251                         cudaMemcpyDeviceToHost));
252
253     transfers_time += wtime() - timing_transfers;
254
255     /* CPU part: Update cluster centers*/
```

```

256
257     timing_cpu = wtime();
258     for (i = 0; i < numObjs; i++) {
259         /* find the array index of nestest cluster center */
260         index = membership[i];
261
262         /* update new cluster centers : sum of objects located within */
263         newClusterSize[index]++;
264         for (j = 0; j < numCoords; j++)
265             newClusters[j][index] += objects[i * numCoords + j];
266     }
267
268     /* average the sum and replace old cluster centers with newClusters */
269     for (i = 0; i < numClusters; i++) {
270         for (j = 0; j < numCoords; j++) {
271             if (newClusterSize[i] > 0)
272                 dimClusters[j][i] = newClusters[j][i] / newClusterSize[i];
273             newClusters[j][i] = 0.0; /* set back to 0 */
274         }
275         newClusterSize[i] = 0; /* set back to 0 */
276     }
277
278     delta /= numObjs;
279     //printf("delta is %f - ", delta);
280     loop++;
281     //printf("completed loop %d\n", loop);
282     cpu_time += wtime() - timing_cpu;
283
284     timing_internal = wtime() - timing_internal;
285     if (timing_internal < timer_min) timer_min = timing_internal;
286     if (timing_internal > timer_max) timer_max = timing_internal;
287 } while (delta > threshold && loop < loop_threshold);
288
289 /*TODO: Update clusters using dimClusters. Be carefull of layout!!!
clusters[numClusters][numCoords] vs dimClusters[numCoords][numClusters] */
290 for (i = 0; i < numClusters; i++) {
291     for (j = 0; j < numCoords; j++) {
292         clusters[i * numCoords + j] = dimClusters[j][i];
293     }
294 }
295
296 timing = wtime() - timing;
297 printf("nloops = %d : total = %lf ms\n\t-> t_loop_avg = %lf ms\n\t-> t_loop_min = %lf
ms\n\t-> t_loop_max = %lf ms\n\t"
298         "-> t_cpu_avg = %lf ms\n\t-> t_gpu_avg = %lf ms\n\t-> t_transfers_avg = %lf
ms\n\n|-----|\n",
299         loop, 1000 * timing, 1000 * timing / loop, 1000 * timer_min, 1000 * timer_max,
300         1000 * cpu_time / loop, 1000 * gpu_time / loop, 1000 * transfers_time / loop);
301
302 char outfile_name[1024] = {0};
303 sprintf(outfile_name, "Execution_logs/silver1-V100_Sz-%lu_Coo-%d_Cl-%d.csv",
304         numObjs * numCoords * sizeof(double) / (1024 * 1024), numCoords, numClusters);
305 FILE *fp = fopen(outfile_name, "a+");

```

```
306 if (!fp) error("Filename %s did not open successfully, no logging performed\n",
307     outfile_name);
308     fprintf(fp, "%s,%d,%lf,%lf,%lf\n", "Shmem", blockSize, timing / loop, timer_min,
309     timer_max);
310     fclose(fp);
311
312     checkCuda(cudaFree(deviceObjects));
313     checkCuda(cudaFree(deviceClusters));
314     checkCuda(cudaFree(deviceMembership));
315
316     free(dimObjects[0]);
317     free(dimObjects);
318     free(dimClusters[0]);
319     free(dimClusters);
320     free(newClusters[0]);
321     free(newClusters);
322     free(newClusterSize);
323
324     return;
325 }
```