a5/cuda_kmeans_transpose.cu

```
#include <stdio.h>
#include <stdlib.h>

#include "kmeans.h"
#include "alloc.h"
#include "error.h"

#ifdef __CUDACC__
inline void checkCuda(cudaError_t e) {
    if (e != cudaSuccess) {
        // cudaGetErrorString() isn't always very helpful. Look up the error
        // number in the cudaError enum in driver_types.h in the CUDA includes
        // directory for a better explanation.
        error("CUDA Error %d: %s\n", e, cudaGetErrorString(e));
    }
}

inline void checkLastCudaError() {
    checkCuda(cudaGetLastError());
}
#endif

__device__ int get_tid() {
  return blockIdx.x * blockDim.x + threadIdx.x;
}

/* square of Euclid distance between two multi-dimensional points using column-base format
*/
__host__ __device__ inline static
double euclid_dist_2_transpose(int numCoords,
                               int numObjs,
                               int numClusters,
                               double *objects,     // [numCoords][numObjs]
                               double *clusters,    // [numCoords][numClusters]
                               int objectId,
                               int clusterId) {
  int i;
  double ans = 0.0;

  /* TODO: Calculate the euclid_dist of elem=objectId of objects from elem=clusterId from
clusters, but for column-base format!!! */
  for (i = 0; i < numCoords; i++) {
    double objectVal = objects[i * numObjs + objectId];
    double clusterVal = clusters[i * numClusters + clusterId];

    double diff = objectVal - clusterVal;
    ans += diff * diff;
  }

  return (ans);
}

```

```
51  __global__ static
52  void find_nearest_cluster(int numCoords,
53                            int numObjs,
54                            int numClusters,
55                            double *objects,           //  [numCoords][numObjs]
56                            double *deviceClusters,    //  [numCoords][numClusters]
57                            int *membership,          //  [numObjs]
58                            double *devdelta) {
59    /* Get the global ID of the thread. */
60    int tid = get_tid();
61
62    if (tid < numObjs) {
63      int index, i;
64      double dist, min_dist;
65
66      /* find the cluster id that has min distance to object */
67      index = 0;
68
69      min_dist = euclid_dist_2_transpose(numCoords, numObjs, numClusters,
70                                objects, deviceClusters,
71                                tid, index);
72
73      for (i = 1; i < numClusters; i++) {
74
75        dist = euclid_dist_2_transpose(numCoords, numObjs, numClusters,
76                              objects, deviceClusters,
77                              tid, i);
78        /* no need square root */
79        if (dist < min_dist) { /* find the min and its array index */
80          min_dist = dist;
81          index = i;
82        }
83      }
84
85      if (membership[tid] != index) {
86
87        atomicAdd(devdelta, 1.0);
88      }
89
90      /* assign the deviceMembership to object objectId */
91      membership[tid] = index;
92    }
93  }
94
95  //
96  //  -----------------------------------------
97  //  DATA LAYOUT
98  //
99  //  objects         [numObjs][numCoords]
100 //  clusters        [numClusters][numCoords]
101 //  dimObjects      [numCoords][numObjs]
102 //  dimClusters     [numCoords][numClusters]
103 //  newClusters     [numCoords][numClusters]
104 //  deviceObjects   [numCoords][numObjs]
```

```
105 //  deviceClusters  [numCoords][numClusters]
106 //  ----------------------------------------
107 //
108 /* return an array of cluster centers of size [numClusters][numCoords]       */
109 void kmeans_gpu(double *objects,      /* in: [numObjs][numCoords] */
110                 int numCoords,    /* no. features */
111                 int numObjs,      /* no. objects */
112                 int numClusters,  /* no. clusters */
113                 double threshold,    /* % objects change membership */
114                 long loop_threshold,   /* maximum number of iterations */
115                 int *membership,   /* out: [numObjs] */
116                 double *clusters,   /* out: [numClusters][numCoords] */
117                 int blockSize) {
118   double timing = wtime(), timing_internal, timer_min = 1e42, timer_max = 0;
119   double timing_gpu, timing_cpu, timing_transfers, transfers_time = 0.0, cpu_time = 0.0,
    gpu_time = 0.0;
120   int loop_iterations = 0;
121   int i, j, index, loop = 0;
122   int *newClusterSize; /* [numClusters]: no. objects assigned in each
123                                   new cluster */
124   double delta = 0, *dev_delta_ptr;          /* % of objects change their clusters */
125
126   /* TODO: Transpose dims */
127   double **dimObjects = (double **) calloc_2d(numCoords, numObjs, sizeof(double));
    //calloc_2d(...) -> [numCoords][numObjs]
128   double **dimClusters = (double **) calloc_2d(numCoords, numClusters, sizeof(double));
    //calloc_2d(...) -> [numCoords][numClusters]
129   double **newClusters = (double **) calloc_2d(numCoords, numClusters, sizeof(double));
    //calloc_2d(...) -> [numCoords][numClusters]
130
131   double *deviceObjects;
132   double *deviceClusters;
133   int *deviceMembership;
134
135   printf("\n|-----------Transpose GPU Kmeans-----------|\n\n");
136
137   //  TODO: Copy objects given in [numObjs][numCoords] layout to new
138   //  [numCoords][numObjs] layout
139   for (i=0 ; i <numObjs; i++){
140     for (j=0; j<numCoords; j++){
141       dimObjects[j][i]=objects[i*numCoords + j];
142     }
143   }
144
145   /* pick first numClusters elements of objects[] as initial cluster centers*/
146   for (i = 0; i < numCoords; i++) {
147     for (j = 0; j < numClusters; j++) {
148       dimClusters[i][j] = dimObjects[i][j];
149     }
150   }
151
152   /* initialize membership[] */
153   for (i = 0; i < numObjs; i++) membership[i] = -1;
154
```

```
155   /* need to initialize newClusterSize and newClusters[0] to all 0 */
156   newClusterSize = (int *) calloc(numClusters, sizeof(int));
157   assert(newClusterSize != NULL);
158
159   timing = wtime() - timing;
160   printf("t_alloc: %lf ms\n\n", 1000 * timing);
161   timing = wtime();
162
163   const unsigned int numThreadsPerClusterBlock = (numObjs > blockSize) ? blockSize :
      numObjs;
164   const unsigned int numClusterBlocks = (numObjs + numThreadsPerClusterBlock - 1) /
      numThreadsPerClusterBlock;
165   const unsigned int clusterBlockSharedDataSize = 0;
166
167   checkCuda(cudaMalloc(&deviceObjects, numObjs * numCoords * sizeof(double)));
168   checkCuda(cudaMalloc(&deviceClusters, numClusters * numCoords * sizeof(double)));
169   checkCuda(cudaMalloc(&deviceMembership, numObjs * sizeof(int)));
170   checkCuda(cudaMalloc(&dev_delta_ptr, sizeof(double)));
171   timing = wtime() - timing;
172   printf("t_alloc_gpu: %lf ms\n\n", 1000 * timing);
173   timing = wtime();
174
175   checkCuda(cudaMemcpy(deviceObjects, dimObjects[0],
176                        numObjs * numCoords * sizeof(double), cudaMemcpyHostToDevice));
177   checkCuda(cudaMemcpy(deviceMembership, membership,
178                        numObjs * sizeof(int), cudaMemcpyHostToDevice));
179   timing = wtime() - timing;
180   printf("t_get_gpu: %lf ms\n\n", 1000 * timing);
181   timing = wtime();
182
183   do {
184     timing_internal = wtime();
185
186     /* GPU part: calculate new memberships */
187
188     timing_transfers = wtime();
189     // TODO: Copy clusters to deviceClusters
190     checkCuda(cudaMemcpy(deviceClusters, dimClusters[0],
191                          numClusters * numCoords * sizeof(double),
192                          cudaMemcpyHostToDevice));
193
194     transfers_time += wtime() - timing_transfers;
195
196     checkCuda(cudaMemset(dev_delta_ptr, 0, sizeof(double)));
197
198     //printf("Launching find_nearest_cluster Kernel with grid_size = %d, block_size = %d,
      shared_mem = %d KB\n", numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDa-
      taSize/1000);
199     timing_gpu = wtime();
200     find_nearest_cluster
201     <<< numClusterBlocks, numThreadsPerClusterBlock, clusterBlockSharedDataSize >>>
202             (numCoords, numObjs, numClusters,
203              deviceObjects, deviceClusters, deviceMembership, dev_delta_ptr);
204
```

```
    cudaDeviceSynchronize();
    checkLastCudaError();
    gpu_time += wtime() - timing_gpu;
    //printf("Kernels complete for itter %d, updating data in CPU\n", loop);

    timing_transfers = wtime();

    checkCuda(cudaMemcpy(membership, deviceMembership,
                    numObjs * sizeof(int),
                    cudaMemcpyDeviceToHost));

    checkCuda(cudaMemcpy(&delta, dev_delta_ptr,
                    sizeof(double),
                    cudaMemcpyDeviceToHost));
    transfers_time += wtime() - timing_transfers;

    /* CPU part: Update cluster centers*/

    timing_cpu = wtime();
    for (i = 0; i < numObjs; i++) {
      /* find the array index of nestest cluster center */
      index = membership[i];

      /* update new cluster centers : sum of objects located within */
      newClusterSize[index]++;
      for (j = 0; j < numCoords; j++)
        newClusters[j][index] += objects[i * numCoords + j];
    }

    /* average the sum and replace old cluster centers with newClusters */
    for (i = 0; i < numClusters; i++) {
      for (j = 0; j < numCoords; j++) {
        if (newClusterSize[i] > 0)
          dimClusters[j][i] = newClusters[j][i] / newClusterSize[i];
        newClusters[j][i] = 0.0;   /* set back to 0 */
      }
      newClusterSize[i] = 0;   /* set back to 0 */
    }

    delta /= numObjs;
    //printf("delta is %f - ", delta);
    loop++;
    //printf("completed loop %d\n", loop);
    cpu_time += wtime() - timing_cpu;

    timing_internal = wtime() - timing_internal;
    if (timing_internal < timer_min) timer_min = timing_internal;
    if (timing_internal > timer_max) timer_max = timing_internal;
  } while (delta > threshold && loop < loop_threshold);

  /*TODO: Update clusters using dimClusters. Be carefull of layout!!!
clusters[numClusters][numCoords] vs dimClusters[numCoords][numClusters] */
  for (i = 0; i < numClusters; i++) {
      for (j = 0; j < numCoords; j++) {
```

```
            clusters[i * numCoords + j] = dimClusters[j][i];
        }
    }

    timing = wtime() - timing;
    printf("nloops = %d  : total = %lf ms\n\t-> t_loop_avg = %lf ms\n\t-> t_loop_min = %lf ms\n\t-> t_loop_max = %lf ms\n\t"
           "-> t_cpu_avg = %lf ms\n\t-> t_gpu_avg = %lf ms\n\t-> t_transfers_avg = %lf ms\n\n|-----------------------------------------|\n",
           loop, 1000 * timing, 1000 * timing / loop, 1000 * timer_min, 1000 * timer_max,
           1000 * cpu_time / loop, 1000 * gpu_time / loop, 1000 * transfers_time / loop);

    char outfile_name[1024] = {0};
    sprintf(outfile_name, "Execution_logs/silver1-V100_Sz-%lu_Coo-%d_Cl-%d.csv",
            numObjs * numCoords * sizeof(double) / (1024 * 1024), numCoords, numClusters);
    FILE *fp = fopen(outfile_name, "a+");
    if (!fp) error("Filename %s did not open succesfully, no logging performed\n", outfile_name);
    fprintf(fp, "%s,%d,%lf,%lf,%lf\n", "Transpose", blockSize, timing / loop, timer_min, timer_max);
    fclose(fp);

    checkCuda(cudaFree(deviceObjects));
    checkCuda(cudaFree(deviceClusters));
    checkCuda(cudaFree(deviceMembership));

    free(dimObjects[0]);
    free(dimObjects);
    free(dimClusters[0]);
    free(dimClusters);
    free(newClusters[0]);
    free(newClusters);
    free(newClusterSize);

    return;
}
```