

```

a6/kmeans/main.c

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>      /* strtok() */
4 #include <sys/types.h>   /* open() */
5 #include <sys/stat.h>
6 #include <fcntl.h>
7 #include <unistd.h>      /* getopt() */
8 #include <mpi.h>
9
10 int _debug;
11 #include "kmeans.h"
12
13 static void usage(char *argv0) {
14     char *help =
15         "Usage: %s [switches]\n"
16         "      -c num_clusters      : number of clusters (must be > 1)\n"
17         "      -s size                : size of examined dataset\n"
18         "      -n num_coords          : number of coordinates\n"
19         "      -t threshold            : threshold value (default : 0.001)\n"
20         "      -l loop_threshold       : iterations threshold (default : 10)\n"
21         "      -d                      : enable debug mode\n"
22         "      -h                      : print this help information";
23     fprintf(stderr, help, argv0);
24     exit(-1);
25 }
26
27 int main(int argc, char **argv)
28 {
29     long i, j, opt;
30     extern char* optarg;
31     extern int optind;
32
33     long    numClusters=0, numCoords=0, numObjs=0;
34     long    rank_numObjs=0;
35     int    * membership; // [rank_numObjs] this array will contain membership
information for this rank's objects
36     int    * tot_membership; // [numObjs]      this array will contain membership
information for all objects
37     double * objects;      // [numObjs * numCoords] data objects
38     double * clusters;     // [numClusters * numCoords] cluster center
39     double  dataset_size = 0, threshold;
40     long    loop_threshold;
41     double  io_timing_read;
42
43     /* some default values */
44     _debug        = 0;
45     threshold     = 0.001;
46     loop_threshold = 10;
47     numClusters   = 0;
48
49     while ( (opt = getopt(argc,argv,"n:t:l:c:s:dh")) != EOF) {
50         switch (opt) {

```

```

51     case 'c': numClusters = atol(optarg);
52             break;
53     case 't': threshold=atof(optarg);
54             break;
55     case 'l': loop_threshold=atol(optarg);
56             break;
57     case 's': dataset_size=atof(optarg);
58             break;
59     case 'n': numCoords=atol(optarg);
60             break;
61     case 'd': _debug = 1;
62             break;
63     case 'h':
64         default: usage(argv[0]);
65             break;
66     }
67 }
68 if (numClusters <= 1) {
69     usage(argv[0]);
70 }
71
72 int rank, size;
73 MPI_Init(&argc,&argv);
74 MPI_Comm_rank(MPI_COMM_WORLD,&rank);
75 MPI_Comm_size(MPI_COMM_WORLD,&size);
76
77 numObjs = (dataset_size*1024*1024) / (numCoords*sizeof(double));
78
79 if (numObjs < numClusters) {
80     if (rank == 0) printf("Error: number of clusters must be larger than the number of
81 data points to be clustered.\n");
82     MPI_Finalize();
83     return 1;
84 }
85 if (rank == 0) printf("dataset_size = %.2f MB      numObjs = %ld      numCoords = %ld
86 numClusters = %ld\n", dataset_size, numObjs, numCoords, numClusters);
87
88 // Allocate space for clusters (coordinates of cluster centers)
89 clusters = (double*) malloc(numClusters * numCoords * sizeof(double));
90
91 // The first numClusters elements are selected as initial centers. Only rank 0 needs
92 // to calculate this, and later broadcast it to all ranks.
93 if (rank == 0) {
94     for (i=0; i<numClusters; i++)
95         for (j=0; j<numCoords; j++)
96             clusters[i*numCoords + j] = objects[i*numCoords + j];
97
98     // check initial cluster centers for repetition
99     if (check_repeated_clusters(numClusters, numCoords, clusters) == 0) {
100         printf("Error: some initial clusters are repeated. Please select distinct
initial centers\n");
101         MPI_Finalize();

```

```
101         return 1;
102     }
103     /*
104     printf("Initial cluster centers:\n");
105     for (i=0; i<numClusters; i++) {
106         printf("(0) clusters[%ld] = ",i);
107         for (j=0; j<numCoords; j++)
108             printf(" %6.6f", clusters[i*numCoords + j]);
109         printf("\n");
110     }
111 */
112 }
113 /*
114 * TODO: Broadcast initial cluster positions to all ranks
115 */
116 MPI_Bcast(clusters, numClusters * numCoords, MPI_DOUBLE, 0, MPI_COMM_WORLD);
117
118
119
120 // membership: the cluster id for each data object
121 membership = (int*) malloc(rank_numObjs * sizeof(int));
122 tot_membership = (int*) malloc(numObjs * sizeof(int));
123
124 // start the core computation
125 /*
126 * TODO: Fix number of objects that this kmeans function call will process
127 */
128 kmeans(objects, numCoords, rank_numObjs, numClusters, threshold, loop_threshold,
membership, clusters);
129
130 /*
131 if (rank == 0) {
132     printf("Final cluster centers:\n");
133     for (i=0; i<numClusters; i++) {
134         printf("clusters[%ld] = ",i);
135         for (j=0; j<numCoords; j++)
136             printf(" %6.6f ", clusters[i*numCoords + j]);
137         printf("\n");
138     }
139 }
140 */
141
142 // Gather membership information from all ranks to tot_membership
143 int recvcounts[size], displs[size];
144 if (rank == 0) {
145     /* TODO: Calculate recvcounts and displs, which will be used to gather data from
each rank.
146         * Hint: recvcounts: number of elements received from each rank
147         *       displs: displacement of each rank's data
148         */
149     int sum_disp = 0;
150     int remainder = numObjs % size;
151
152     for (j = 0; j < size; j++) {
```

```
153     // Υπολογισμός πόσα αντικείμενα περιμένουμε από το rank 'j'  
154     int j_numObjs = numObjs / size;  
155     if (j < remainder) {  
156         j_numObjs++;  
157     }  
158  
159     recvcounts[j] = j_numObjs; // Εδώ είναι σκέτα αντικείμενα (int)  
160     displs[j] = sum_disp;  
161     sum_disp += recvcounts[j];  
162 }  
163 }  
164  
165 /*  
166 * TODO: Broadcast the recvcounts and displs arrays to other ranks.  
167 */  
168 MPI_Bcast(recvcounts, size, MPI_INT, 0, MPI_COMM_WORLD);  
169 MPI_Bcast(displs, size, MPI_INT, 0, MPI_COMM_WORLD);  
170  
171 /*  
172 * TODO: Gather membership information from every rank. (hint: each rank may send  
173 different number of objects)  
174 */  
175 MPI_Gatherv(membership, rank_numObjs, MPI_INT, tot_membership, recvcounts, displs,  
MPI_INT, 0, MPI_COMM_WORLD);  
176  
177  
178 if (_debug && rank == 0)  
179     for (i = 0; i < numObjs; ++i)  
180         fprintf(stderr, "%d\n", tot_membership[i]);  
181  
182 free(objects);  
183 free(membership);  
184 free(tot_membership);  
185 free(clusters);  
186  
187 MPI_Finalize();  
188 return 0;  
189 }  
190 }
```